

Course Name: Deep Learning

Lab Title: Vehicle Detection for Smart Traffic Management using YOLOv11

Student Name:Manavi Pawar

PRN:-202201040050

Date of Submission: 01-04-2025

Group Members: Manavi Pawar(202201040050),Kanishka Garud(202201070062),Sakshi Dube(202201040155)

Objective The purpose of this lab is to understand and implement YOLOv11 for real-time object detection. Students will perform dataset preparation, model implementation, inference, and performance evaluation.

Start coding or [generate](#) with AI.

Task 1: Environment Setup and YOLOv11 Installation

Objective: Set up the required libraries and dependencies to run YOLOv11.

Instructions:

Install Python and required libraries (PyTorch, OpenCV, Ultralytics, etc.). Install YOLOv11 from the official repository. Verify the installation by running a sample script. Expected Outcome: A functional YOLOv11 environment ready for experimentation.

```
!pip install ultralytics
```



```
Attempting uninstall: nvidia-cusolver-cu12
Found existing installation: nvidia-cusolver-cu12 11.6.3.83
Uninstalling nvidia-cusolver-cu12-11.6.3.83:
Successfully uninstalled nvidia-cusolver-cu12-11.6.3.83
Successfully installed nvidia-cublas-cu12-12.4.5.8 nvidia-cuda-cupti-cu12-12.4.127 nvidia-cuda-nvrtc-cu12-12.4.127 nvidia-cud
```

```
from ultralytics import YOLO

# Load a pre-trained YOLOv11 model
model = YOLO('yolo11n.pt') # 'n' stands for nano version; other versions include 's', 'm', 'l', 'x'

# Run YOLO on a sample image
results = model('https://ultralytics.com/images/zidane.jpg')
results[0].show()
```

🔗 Creating new Ultralytics Settings v0.0.6 file ☒
View Ultralytics Settings with 'yolo settings' or at '/root/.config/Ultralytics/settings.json'
Update Settings with 'yolo settings key=value', i.e. 'yolo settings runs_dir=path/to/dir'. For help see <https://docs.ultralytics.com>
Downloading <https://github.com/ultralytics/assets/releases/download/v8.3.0/yolo11n.pt> to 'yolo11n.pt'...
100%|██████████| 5.35M/5.35M [00:00<00:00, 62.1MB/s]

Downloading <https://ultralytics.com/images/zidane.jpg> to 'zidane.jpg'...
100%|██████████| 49.2k/49.2k [00:00<00:00, 4.73MB/s]
image 1/1 /content/zidane.jpg: 384x640 2 persons, 1 tie, 364.4ms
Speed: 18.1ms preprocess, 364.4ms inference, 38.9ms postprocess per image at shape (1, 3, 384, 640)



Task 2: Dataset Preparation & Preprocessing Objective: Load and preprocess a dataset for object detection.

Instructions:

Choose a Dataset – Use COCO, Pascal VOC, or a custom dataset. Annotate Images – If using a custom dataset, label objects using Roboflow or Labellmg. Convert Annotations – Use Roboflow to export the dataset in YOLO format. Download the Dataset – Use the Roboflow API to fetch the dataset. Split the Dataset – Divide into train (80%), validation (10%), and test (10%). Expected Outcome: A well-structured dataset in YOLO format.

```
!pip install roboflow
```

```
➦ Requirement already satisfied: roboflow in /usr/local/lib/python3.11/dist-packages (1.1.58)
Requirement already satisfied: certifi in /usr/local/lib/python3.11/dist-packages (from roboflow) (2025.1.31)
Requirement already satisfied: idna==3.7 in /usr/local/lib/python3.11/dist-packages (from roboflow) (3.7)
Requirement already satisfied: cyclor in /usr/local/lib/python3.11/dist-packages (from roboflow) (0.12.1)
Requirement already satisfied: kiwisolver>=1.3.1 in /usr/local/lib/python3.11/dist-packages (from roboflow) (1.4.8)
Requirement already satisfied: matplotlib in /usr/local/lib/python3.11/dist-packages (from roboflow) (3.10.0)
Requirement already satisfied: numpy>=1.18.5 in /usr/local/lib/python3.11/dist-packages (from roboflow) (2.0.2)
Requirement already satisfied: opencv-python-headless==4.10.0.84 in /usr/local/lib/python3.11/dist-packages (from roboflow) (4.10.0.84)
Requirement already satisfied: Pillow>=7.1.2 in /usr/local/lib/python3.11/dist-packages (from roboflow) (11.1.0)
Requirement already satisfied: pillow-heif>=0.18.0 in /usr/local/lib/python3.11/dist-packages (from roboflow) (0.22.0)
Requirement already satisfied: python-dateutil in /usr/local/lib/python3.11/dist-packages (from roboflow) (2.8.2)
Requirement already satisfied: python-dotenv in /usr/local/lib/python3.11/dist-packages (from roboflow) (1.1.0)
Requirement already satisfied: requests in /usr/local/lib/python3.11/dist-packages (from roboflow) (2.32.3)
Requirement already satisfied: six in /usr/local/lib/python3.11/dist-packages (from roboflow) (1.17.0)
Requirement already satisfied: urllib3>=1.26.6 in /usr/local/lib/python3.11/dist-packages (from roboflow) (2.3.0)
Requirement already satisfied: tqdm>=4.41.0 in /usr/local/lib/python3.11/dist-packages (from roboflow) (4.67.1)
Requirement already satisfied: PyYAML>=5.3.1 in /usr/local/lib/python3.11/dist-packages (from roboflow) (6.0.2)
Requirement already satisfied: requests-toolbelt in /usr/local/lib/python3.11/dist-packages (from roboflow) (1.0.0)
Requirement already satisfied: filetype in /usr/local/lib/python3.11/dist-packages (from roboflow) (1.2.0)
Requirement already satisfied: contourpy>=1.0.1 in /usr/local/lib/python3.11/dist-packages (from matplotlib->roboflow) (1.3.1)
Requirement already satisfied: fonttools>=4.22.0 in /usr/local/lib/python3.11/dist-packages (from matplotlib->roboflow) (4.56.0)
Requirement already satisfied: packaging>=20.0 in /usr/local/lib/python3.11/dist-packages (from matplotlib->roboflow) (24.2)
Requirement already satisfied: pyparsing>=2.3.1 in /usr/local/lib/python3.11/dist-packages (from matplotlib->roboflow) (3.2.1)
Requirement already satisfied: charset-normalizer<4,>=2 in /usr/local/lib/python3.11/dist-packages (from requests->roboflow) (3.3.2)
```

```
from roboflow import Roboflow
```

```
rf = Roboflow(api_key="Z36dR2kczFG0Re8fzTeh")
```

```
project = rf.workspace("roboflow-100").project("vehicles-q0x2v")
```

```
➦ loading Roboflow workspace...
loading Roboflow project...
```

```
dataset = project.version(1).download("yolov11")
```

```
➦ Downloading Dataset Version Zip in vehicles-1 to yolov11:: 100%|██████████| 217401/217401 [00:04<00:00, 47393.81it/s]
Extracting Dataset Version Zip to vehicles-1 in yolov11:: 100%|██████████| 8128/8128 [00:04<00:00, 1933.94it/s]
```

```
import os
dataset_path = "/content/vehicles-1"
print(os.listdir(dataset_path))
```

```
➦ ['train', 'data.yaml', 'valid', 'test', 'README.roboflow.txt', 'README.dataset.txt']
```

```
yaml_path = "/content/vehicles-1/data.yaml"
```

```
with open(yaml_path, "r") as file:
    print(file.read())
```

```
➦ train: ../train/images
val: ../valid/images
test: ../test/images

nc: 12
names: ['big bus', 'big truck', 'bus-l-', 'bus-s-', 'car', 'mid truck', 'small bus', 'small truck', 'truck-l-', 'truck-m-', 'truck']

roboflow:
  workspace: roboflow-100
  project: vehicles-q0x2v
  version: 1
  license: CC BY 4.0
  url: https://universe.roboflow.com/roboflow-100/vehicles-q0x2v/dataset/1
```

```
import yaml
```

```
# Load and update YAML file with correct paths
with open(yaml_path, "r") as file:
    data = yaml.safe_load(file)
```

```
correct_paths = {
    "train": "/content/vehicles-1/train",
    "val": "/content/vehicles-1/valid",
    "test": "/content/vehicles-1/test"
```

```
}

data.update(correct_paths)

with open(yaml_path, "w") as file:
    yaml.dump(data, file, default_flow_style=False)

print("✅ data.yaml paths have been updated successfully!")
```

```
➦ ✅ data.yaml paths have been updated successfully!
```

```
with open(yaml_path, "r") as file:
    print(file.read())

➦ train: ../train/images
    val: ../valid/images
    test: ../test/images

    nc: 12
    names: ['big bus', 'big truck', 'bus-l-', 'bus-s-', 'car', 'mid truck', 'small bus', 'small truck', 'truck-l-', 'truck-m-', 'tru

    roboflow:
        workspace: roboflow-100
        project: vehicles-q0x2v
        version: 1
        license: CC BY 4.0
        url: https://universe.roboflow.com/roboflow-100/vehicles-q0x2v/dataset/1
```

```
# Check GPU availability
import torch
print(torch.cuda.is_available()) # Should print True if GPU is available
print(torch.cuda.device_count()) # Number of GPUs available
print(torch.cuda.get_device_name(0) if torch.cuda.is_available() else "No GPU detected")
```

```
➦ False
    0
    No GPU detected
```

Task 3: Training YOLOv11 Model Objective: Train YOLOv11 on the prepared dataset.

Instructions:

Configure the training parameters (batch size, epochs, learning rate). Train the YOLOv11 model using the dataset. Monitor training progress (loss, accuracy, mAP). Save the trained model weights.

Expected Outcome: A trained YOLOv11 model ready for inference.

```
from ultralytics import YOLO
```

```
model = YOLO("yolo11n.pt")
results = model.train(
    data="/content/vehicles-1/data.yaml",
    epochs=5,
    batch=8,
    device='cpu' # Use GPU
)
```

```
➦
```


Epoch	GPU_mem	box_loss	cls_loss	dfl_loss	Instances	Size
3/5	0G	1.323	1.42	1.089	34	640: 100% ██████████ 330/330 [37:49<00:00, 6.88s/
	Class	Images	Instances	Box(P	R	mAP50 mAP50-95): 100% ██████████ 61/61 [05:53<00:00
	all	966	13450	0.473	0.361	0.313 0.204
Epoch	GPU_mem	box_loss	cls_loss	dfl_loss	Instances	Size
4/5	0G	1.296	1.322	1.082	42	640: 100% ██████████ 330/330 [38:21<00:00, 6.97s/
	Class	Images	Instances	Box(P	R	mAP50 mAP50-95): 100% ██████████ 61/61 [05:40<00:00
	all	966	13450	0.427	0.42	0.346 0.229
Epoch	GPU_mem	box_loss	cls_loss	dfl_loss	Instances	Size
5/5	0G	1.263	1.227	1.066	66	640: 100% ██████████ 330/330 [37:34<00:00, 6.83s/
	Class	Images	Instances	Box(P	R	mAP50 mAP50-95): 100% ██████████ 61/61 [05:33<00:00
	all	966	13450	0.443	0.453	0.349 0.234

5 epochs completed in 3.654 hours.
Optimizer stripped from runs/detect/train/weights/last.pt, 5.5MB
Optimizer stripped from runs/detect/train/weights/best.pt, 5.5MB

Validating runs/detect/train/weights/best.pt...
Ultralytics 8.3.96 🚀 Python-3.11.11 torch-2.6.0+cu124 CPU (Intel Xeon 2.20GHz)
YOLO11n summary (fused): 100 layers, 2,584,492 parameters, 0 gradients, 6.3 GFLOPs

Class	Images	Instances	Box(P	R	mAP50	mAP50-95):
all	966	13450	0.443	0.453	0.349	0.233
big bus	210	273	0.705	0.608	0.701	0.521
big truck	404	1162	0.69	0.42	0.588	0.35
bus-l-	8	8	0.0143	0.5	0.0111	0.00494
bus-s-	12	12	1	0	0.0031	0.00257
car	927	8537	0.799	0.772	0.819	0.46
mid truck	118	257	0.399	0.0934	0.126	0.097
small bus	43	49	0	0	0.01	0.0072
small truck	517	1721	0.629	0.513	0.574	0.338
truck-l-	266	433	0.309	0.617	0.38	0.285
truck-m-	331	629	0.271	0.781	0.38	0.283
truck-s-	147	221	0.175	0.498	0.186	0.132
truck-xl-	110	148	0.325	0.628	0.411	0.32

Speed: 5.7ms preprocess, 259.2ms inference, 0.0ms loss, 4.7ms postprocess per image
Results saved to **runs/detect/train**

```
# Save the trained model weights
best_model_path = "/content/runs/detect/train/weights/best.pt"
print(f"Model training complete. Best model saved at: {best_model_path}")
```

🔄 Model training complete. Best model saved at: /content/runs/detect/train/weights/best.pt

Task 4: Model Inference and Evaluation

```
# Load the trained model
model = YOLO(best_model_path)
```

```
image_path = "/content/vehicles-1/test/images/adit_mp4-815_jpg.rf.fb532f30f712174b620afee0cfb1bfbb.jpg"
results = model(image_path, save=True, conf=0.5)
```

🔄 image 1/1 /content/vehicles-1/test/images/adit_mp4-815_jpg.rf.fb532f30f712174b620afee0cfb1bfbb.jpg: 640x640 3 cars, 1 truck-m-, Speed: 5.6ms preprocess, 376.4ms inference, 7.0ms postprocess per image at shape (1, 3, 640, 640)
Results saved to **runs/detect/predict**

```
for result in results:
    result.show()
```



```
# Evaluate the model performance
metrics = model.val()

map_50 = metrics.box.map50 # mAP@50
map_50_95 = metrics.box.map # mAP@50-95
precision = metrics.box.p.mean().item() if metrics.box.p.size > 0 else 0.0
recall = metrics.box.r.mean().item() if metrics.box.r.size > 0 else 0.0

print(f"\U0001F4CA mAP@50: {map_50:.4f}")
print(f"\U0001F4CA mAP@50-95: {map_50_95:.4f}")
print(f"\U0001F4C8 Precision: {precision:.4f}")
print(f"\U0001F4C9 Recall: {recall:.4f}")

if precision + recall > 0:
    f1_score = 2 * (precision * recall) / (precision + recall)
    print(f"\U0001F525 F1 Score: {f1_score:.4f}")
else:
    print("\U0001F6A8 Cannot compute F1 Score (Precision + Recall = 0)")
```

```
Ultralytics 8.3.96 Python-3.11.11 torch-2.6.0+cu124 CPU (Intel Xeon 2.20GHz)
val: Scanning /content/vehicles-1/valid/labels.cache... 966 images, 3 backgrounds, 0 corrupt: 100%|██████████| 966/966 [00:00<?,
Class      Images  Instances  Box(P  R      mAP50  mAP50-95): 100%|██████████| 61/61 [04:42<00:00,
  all        966    13450    0.443  0.453   0.349   0.233
  big bus    210     273     0.705  0.608   0.701   0.521
  big truck  404    1162     0.69   0.42   0.588   0.35
  bus-l-      8        8    0.0143  0.5    0.0111  0.00494
  bus-s-     12       12      1      0    0.0031  0.00257
  car       927    8537     0.799  0.772   0.819   0.46
  mid truck  118     257     0.399  0.0934  0.126   0.097
  small bus  43      49      0      0     0.01    0.0072
  small truck 517    1721     0.629  0.513   0.574   0.338
  truck-l-   266    433     0.309  0.617   0.38    0.285
  truck-m-   331    629     0.271  0.781   0.38    0.283
  truck-s-   147    221     0.175  0.498   0.186   0.132
  truck-xl-  110    148     0.325  0.628   0.411   0.32

Speed: 4.9ms preprocess, 272.9ms inference, 0.0ms loss, 4.7ms postprocess per image
Results saved to runs/detect/val4
mAP@50: 0.3491
mAP@50-95: 0.2334
Precision: 0.4430
Recall: 0.4525
F1 Score: 0.4477
```



```
for result in results:  
    result.show()
```



Conclusion:-

The project "Vehicle Detection for Smart Traffic Management using YOLOv11" successfully implemented real-time object detection to identify vehicles such as cars, buses, and auto-rickshaws. The model demonstrated strong precision, indicating high confidence in its predictions. However, recall was relatively lower, highlighting missed detections.

Challenges were observed in detecting specific violations like helmet and no-helmet detection, which suggests the need for additional training data and optimized hyperparameters. Future improvements could focus on data augmentation, fine-tuning the YOLOv11 model, and exploring alternative architectures to enhance overall detection accuracy.

This project provides a foundational framework for automated traffic monitoring and violation detection, which can be further refined for real-world deployment in smart city applications.

Declaration

I, Manavi Pawar, confirm that the work submitted in this assignment is my own and has been completed following academic integrity guidelines.

Signature: Manavi Pawar